AWS Auto Scaling with Application Load Balancer (ALB)

Step 1: Host a Simple PHP Application in EC2

1. Launch an EC2 instance (Amazon Linux 2 preferred).
2. Connect via SSH:
   ssh -i test.pem ec2-user@65.2.6.16
3. Install Apache and PHP:
   sudo yum install -y httpd php
4. Start and enable Apache:
   sudo systemctl start httpd
   sudo systemctl enable httpd
5. Create a PHP test page:
   echo "<?php phpinfo(); ?>" | sudo tee /var/www/html/index.php
6. Open HTTP (port 80) in the instance's Security Group.
7. Test by visiting http://65.2.6.16/ in a browser.

Step 2: Configure an Application Load Balancer (ALB)

1. Go to EC2 → Load Balancers → Create Load Balancer → Application Load Balancer.
2. Select scheme: Internet-facing, and IP type: IPv4.
3. Choose at least two subnets from different Availability Zones.
4. Select a Security Group allowing inbound HTTP (port 80).
5. Create a Target Group:
   - Target type: Instance
   - Protocol: HTTP
   - Port: 80
   - Health check path: /
6. Register your running EC2 instance with the target group.
7. Finish ALB setup and note the DNS name (e.g., my-alb-12345.ap-south-

1.elb.amazonaws.com).

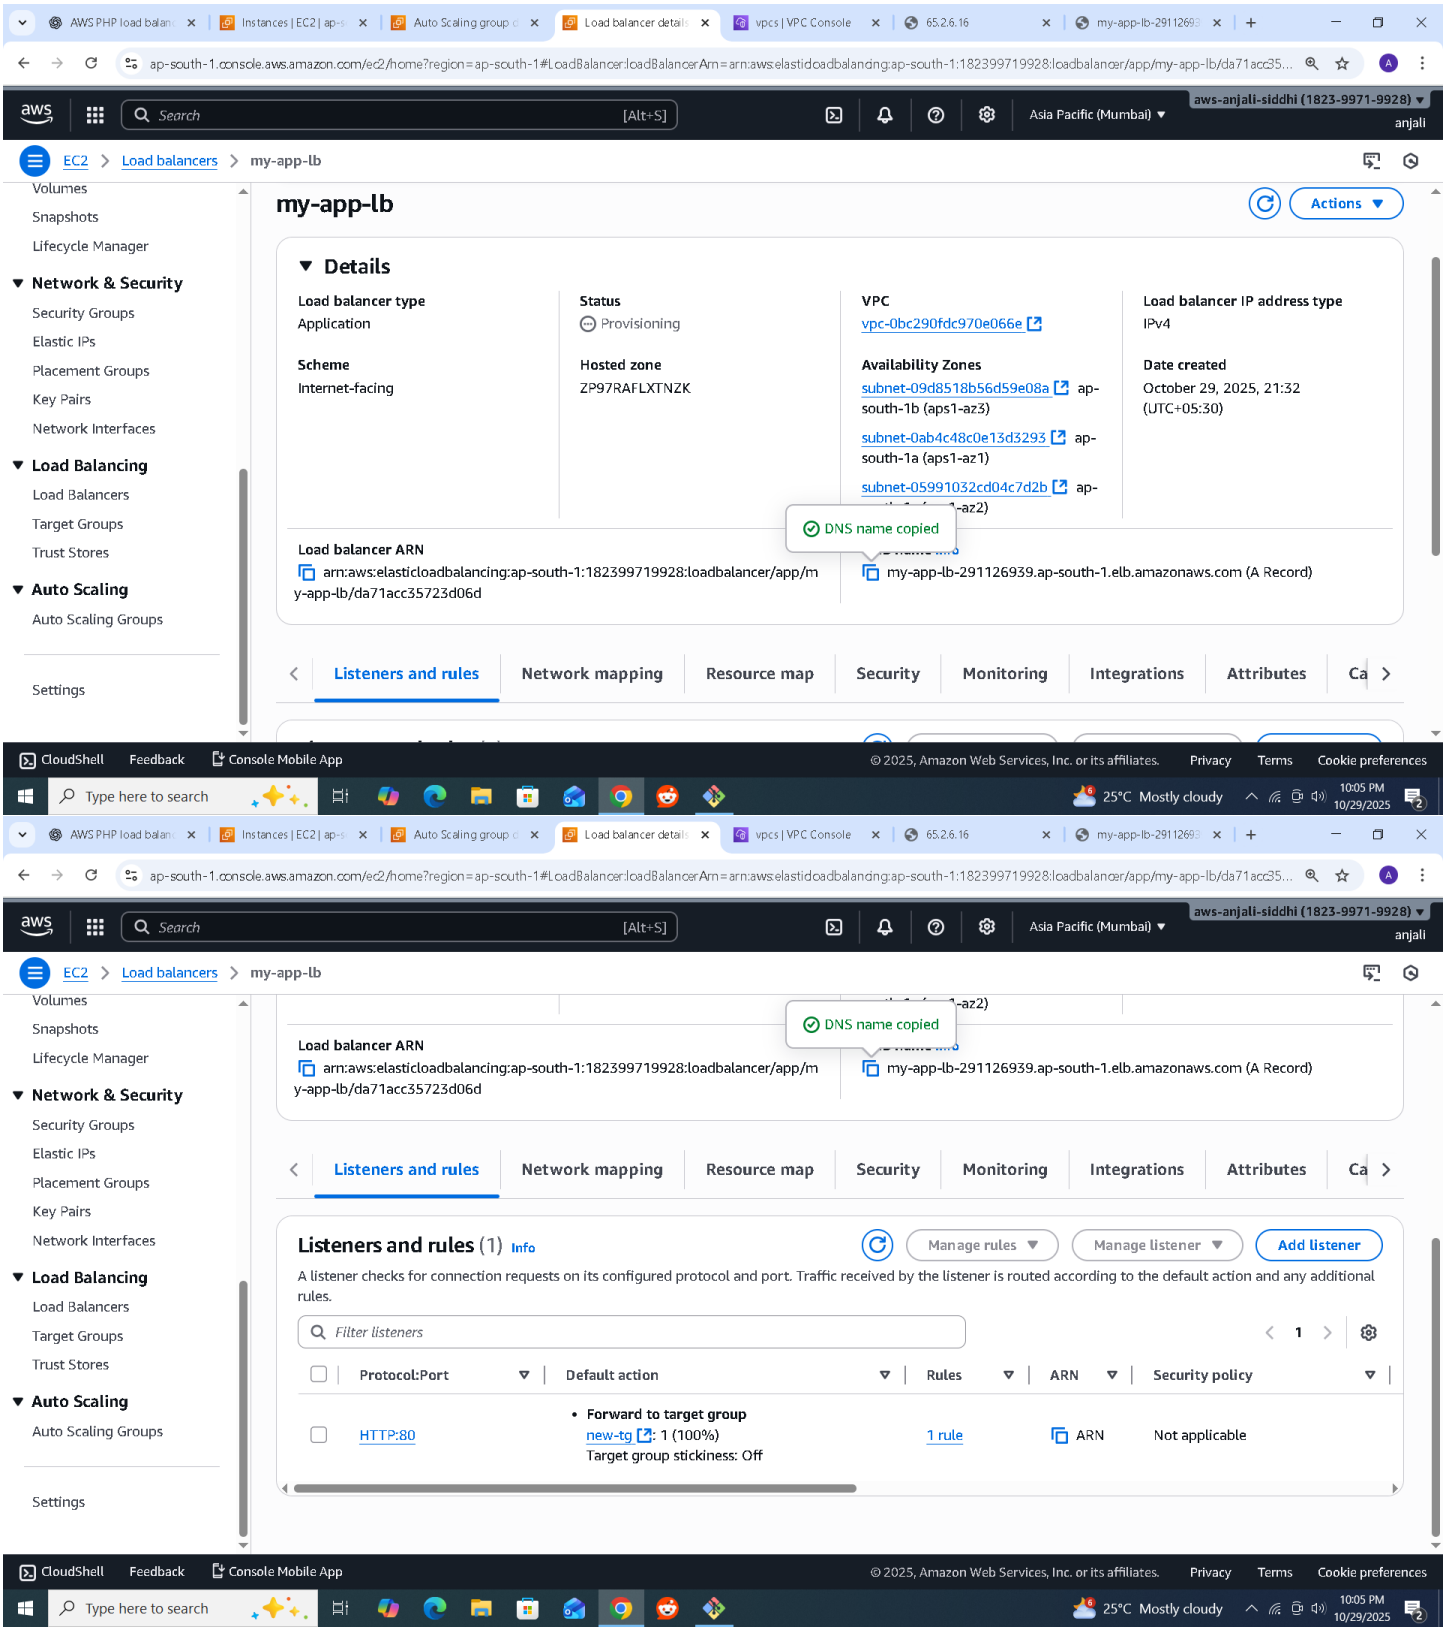

Step 3: Create Launch Template for Auto Scaling

1. Go to EC2 → Launch Templates → Create Launch Template.
2. Configure:
   - Name: auto-scale
   - AMI: same as your instance
   - Instance type: t2.micro
   - Key pair: test
   - Security Group: same as your instance (allow port 80)
3. Leave subnet blank (ASG will handle multi-AZ placement).
4. Save the template.

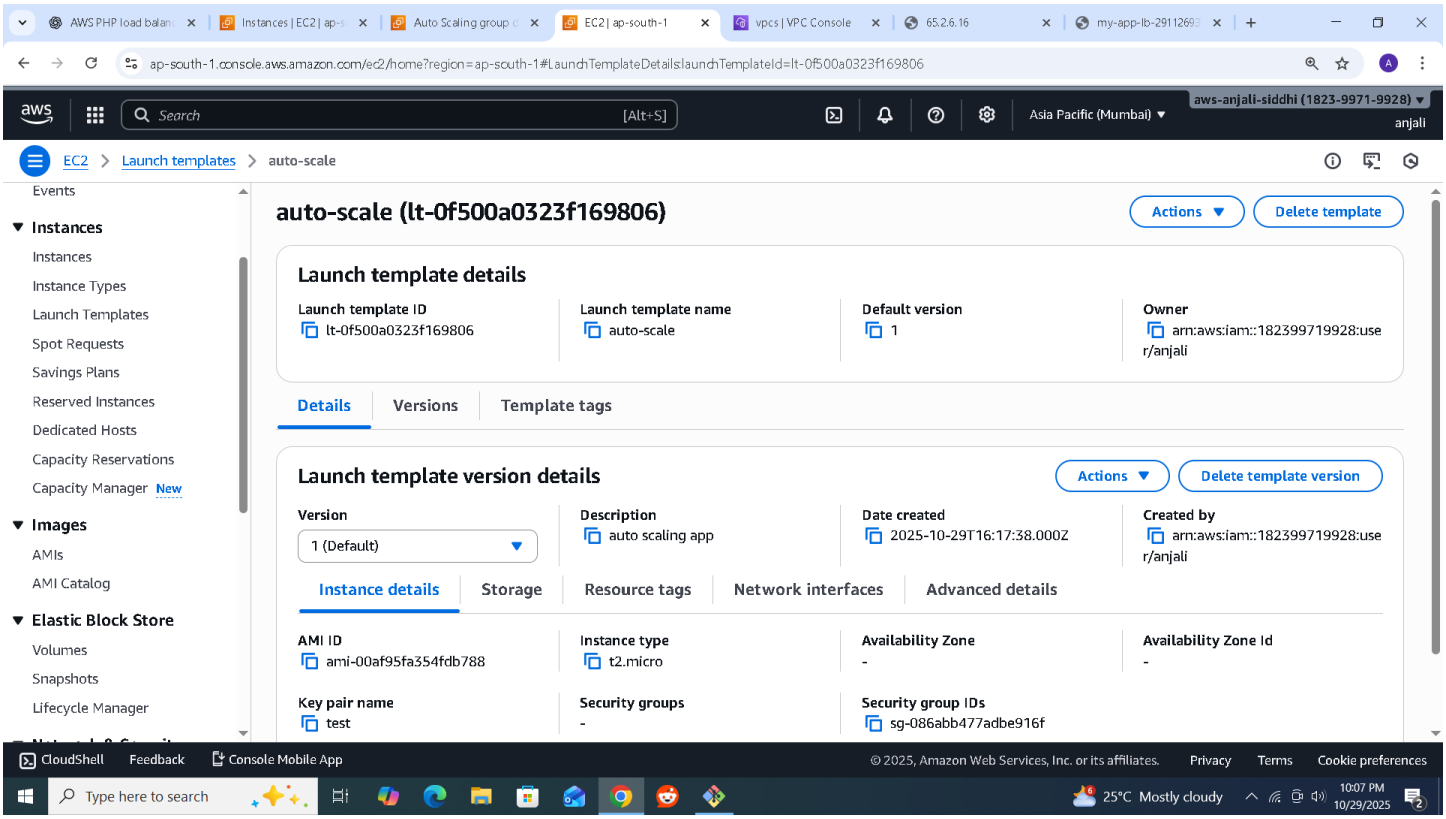

Step 4: Create an Auto Scaling Group (ASG)

1. Go to EC2 → Auto Scaling Groups → Create Auto Scaling Group.
2. Select the Launch Template created earlier.
3. Set:
   - Desired capacity: 1
   - Minimum capacity: 1

- Maximum capacity: 3

4. Choose the same VPC and subnets as the ALB.
5. Attach the previously created Target Group.
6. Enable EC2 and ELB health checks (recommended).
7. Set Health check grace period: 300 seconds.
8. Skip or disable VPC Lattice integration unless required.
9. Finish creation.

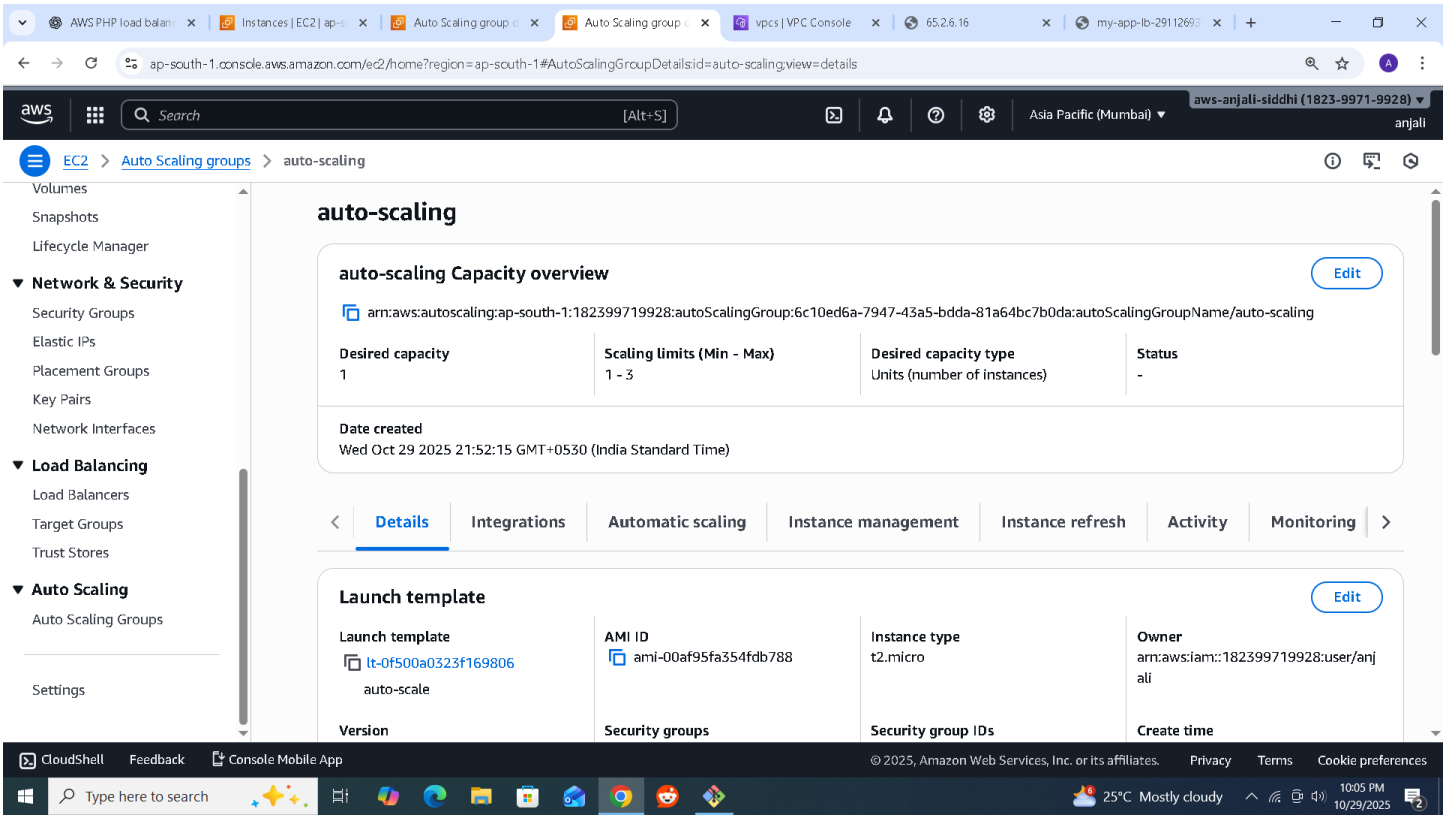

Step 5: Configure Scaling Policy (CPU-based)

1. Go to your Auto Scaling Group → Automatic Scaling tab.
2. Click “Add policy” → Choose “Target tracking scaling policy.”
3. Configure as:
   - Policy name: CPU50-TargetTracking
   - Metric type: Average CPU utilization
   - Target value: 50%
   - Cooldown period: 300 seconds (default)
4. Save the policy.

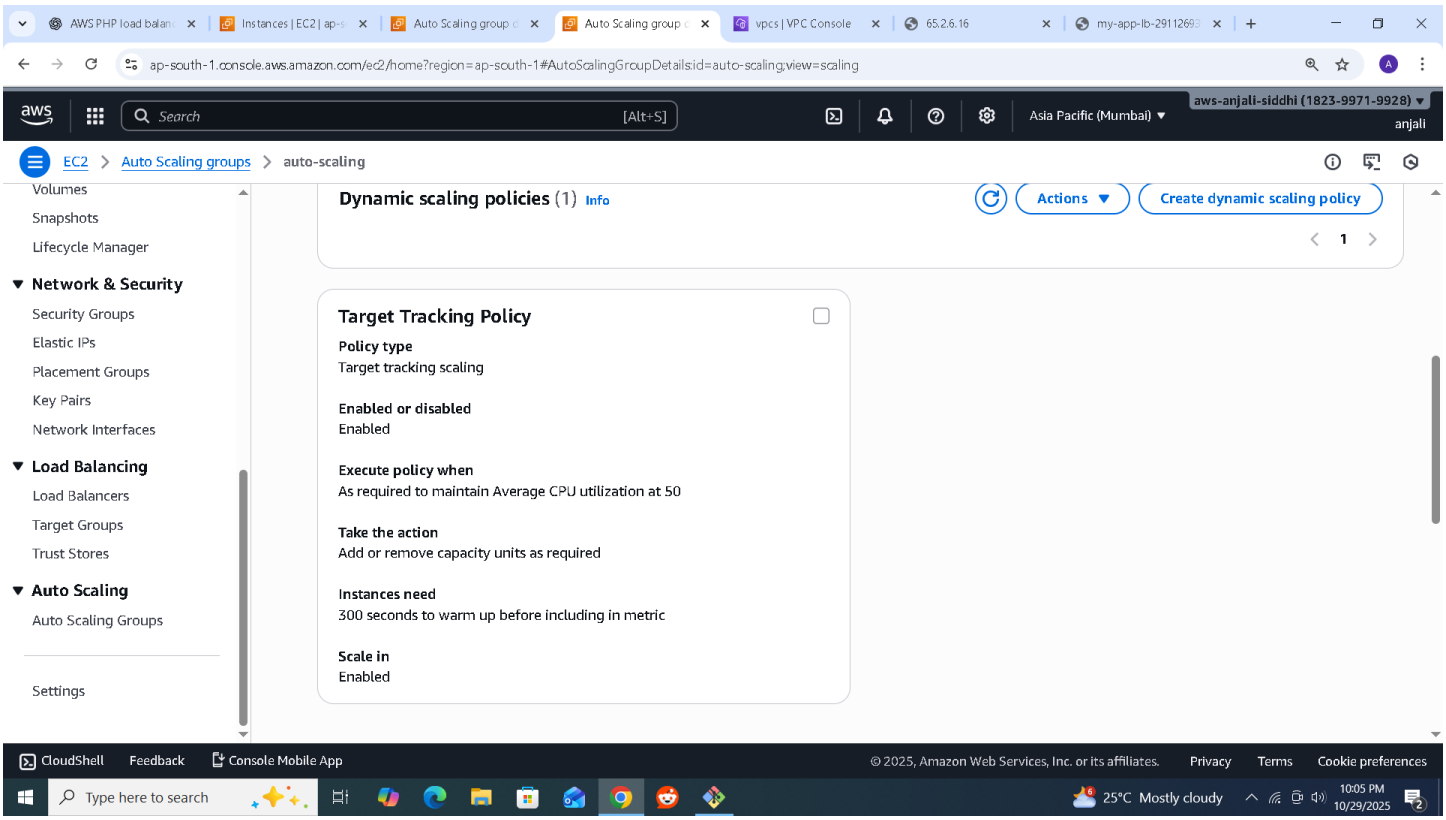

Step 6: Test Scaling Behavior

1. SSH into your instance:
   ssh -i test.pem ec2-user@<instance-public-ip>
2. Install the stress tool:
   sudo yum install -y stress
3. Run the load test:
   stress --cpu 2 --timeout 300
4. Monitor scaling:
   - In EC2 → Auto Scaling Group → Activity tab, watch for new instance launches.
   - In EC2 → Instances, new instances will appear.
   - In Load Balancer → Target Groups, all instances will show as healthy.
5. Once CPU usage drops below 50%, the ASG will terminate extra instances.

Step 7: Verification and Monitoring

1. Verify in CloudWatch:
   - Go to CloudWatch → Metrics → EC2 → CPUUtilization.
   - Observe the spike during the stress test and the automatic scaling response.
2. Visit your ALB DNS URL and ensure the load is distributed across instances.

✓Successfully configured AWS Auto Scaling with an Application Load Balancer for a PHP web application.

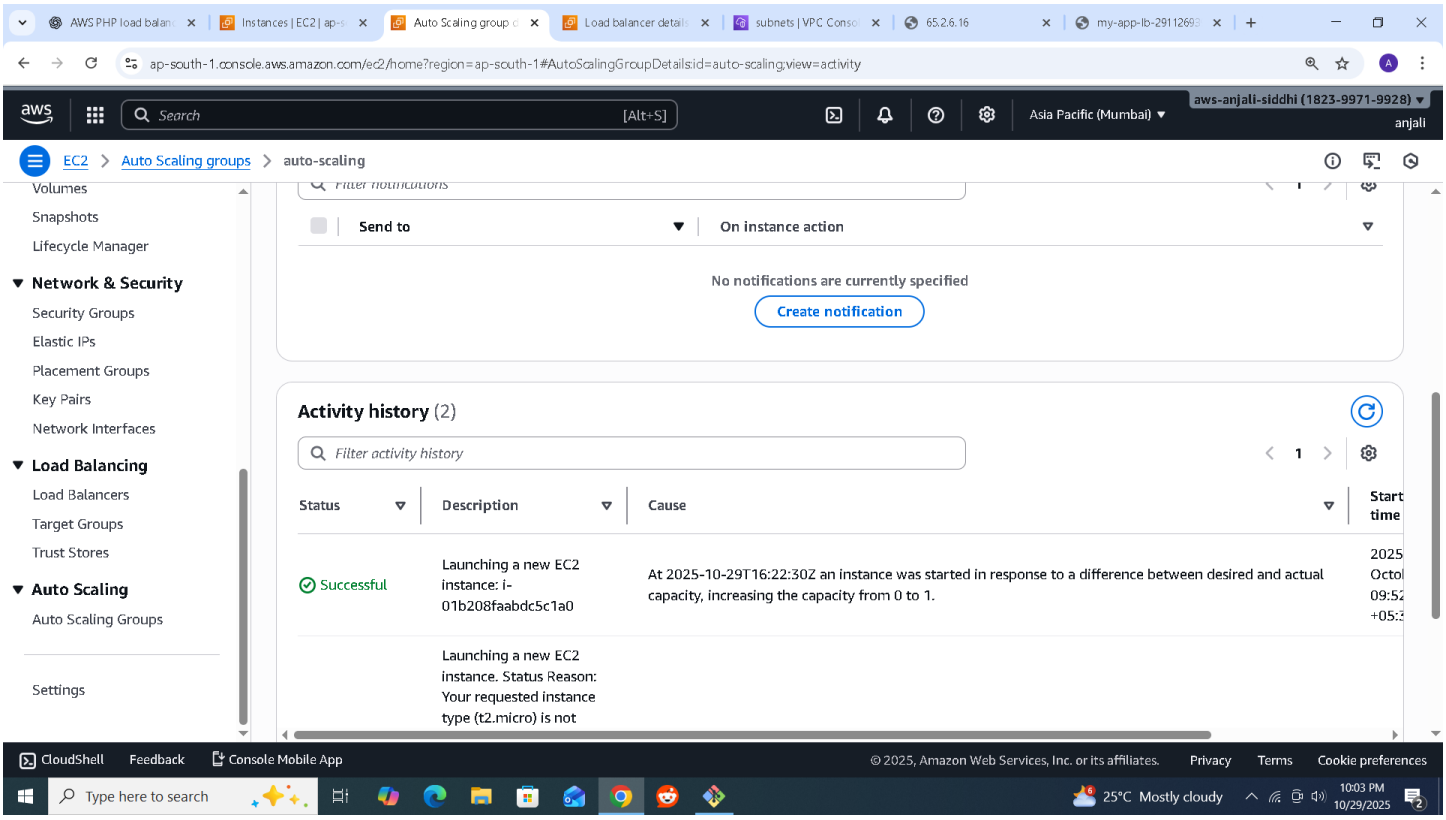

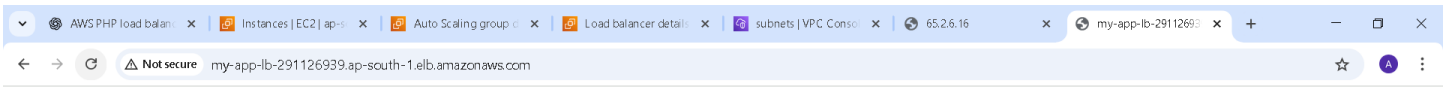

Hello from EC2 instance: ip-172-31-12-61.ap-south-1.compute.internal

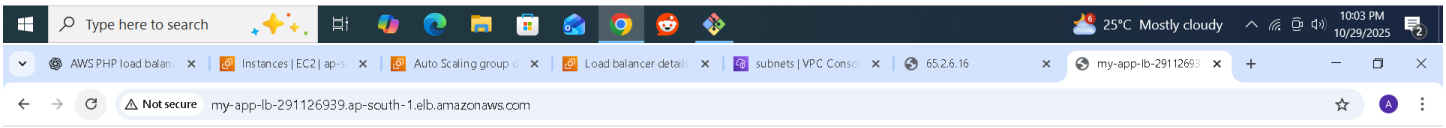

Hello from Auto-Scaled instance: ip-172-31-37-121.ap-south-1.compute.internal

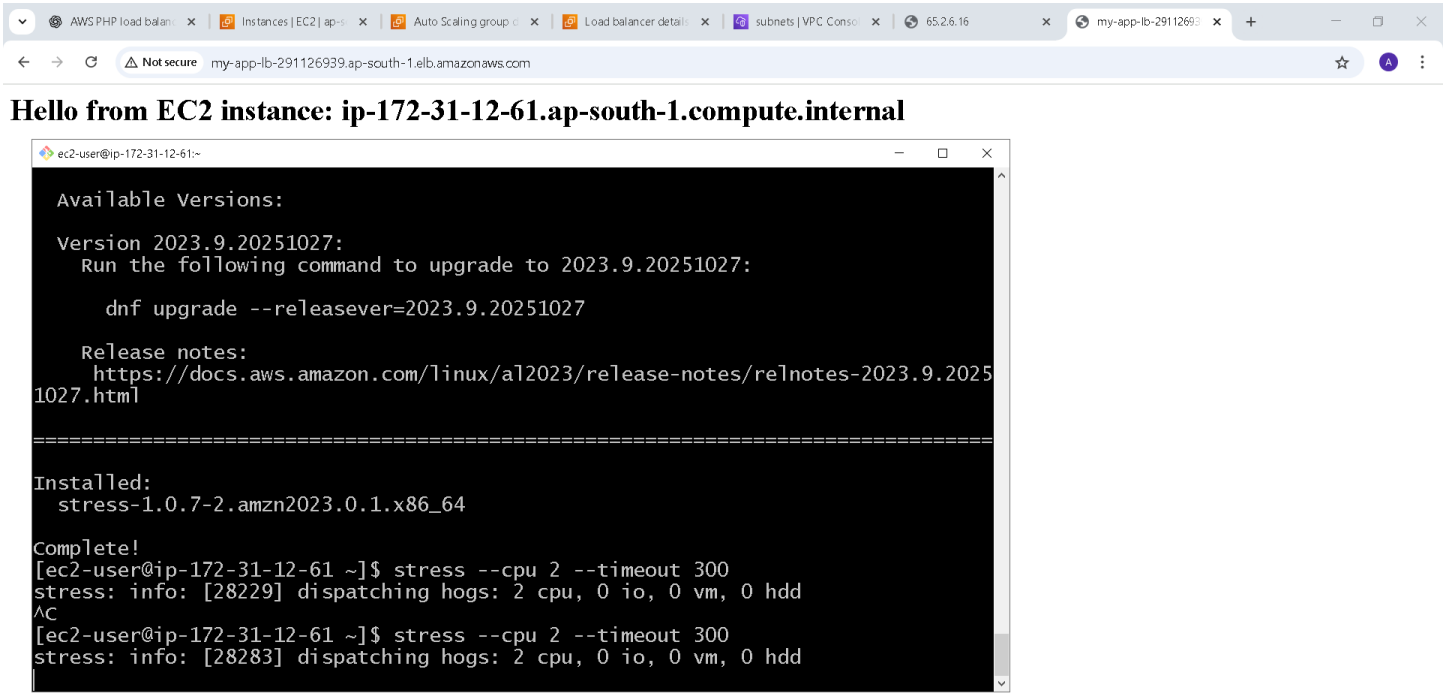
my-app-lb-291126939.ap-south-1.elb.amazonaws.com
Hello from EC2 instance: ip-172-31-12-61.ap-south-1.compute.internal
ec2-user@ip-172-31-12-61:~
Available Versions:
Version 2023.9.20251027:
Run the following command to upgrade to 2023.9.20251027:
dnf upgrade --releasever=2023.9.20251027
Release notes:
https://docs.aws.amazon.com/linux/al2023/release-notes/relnotes-2023.9.20251027.html
Installed:
stress-1.0.7-2.amzn2023.0.1.x86_64
Complete!
[ec2-user@ip-172-31-12-61 ~]$ stress --cpu 2 --timeout 300
stress: info: [28229] dispatching hogs: 2 cpu, 0 io, 0 vm, 0 hdd
^C
[ec2-user@ip-172-31-12-61 ~]$ stress --cpu 2 --timeout 300
stress: info: [28283] dispatching hogs: 2 cpu, 0 io, 0 vm, 0 hdd

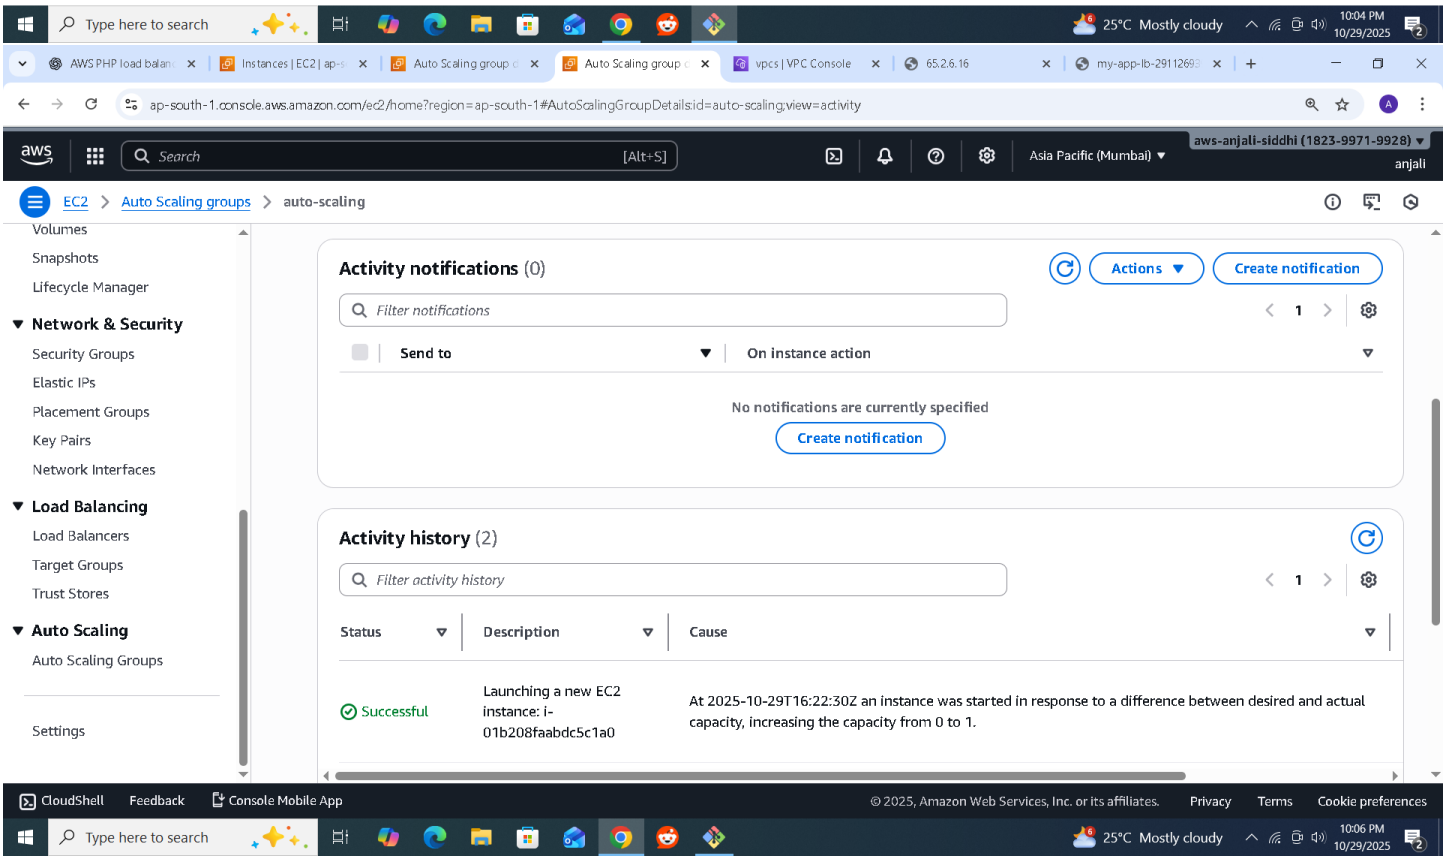
ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#AutoScalingGroupDetails:id=auto-scaling;view=activity
EC2 > Auto Scaling groups > auto-scaling
Activity notifications (0)
Actions
Create notification
Filter notifications
Send to
On instance action
No notifications are currently specified
Create notification
Activity history (2)
Filter activity history
Status
Description
Cause
Successful
Launching a new EC2 instance: i-01b208faabdc5c1a0
At 2025-10-29T16:22:30Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1.
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Trust Stores
Auto Scaling
Auto Scaling Groups
Settings
CloudShell
Feedback
Console Mobile App
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences